

# Class 7: Machine Learning 1

Melanie Alonzo (A17375327)

## Table of contents

Background . . . . .	1
K-means clustering . . . . .	3
Hierarchical CLustering . . . . .	7
PCA of UK food data . . . . .	9
Spotting manjor differences and trends . . . . .	11
Paris Plots and heat maps . . . . .	13
PCA to the rescue . . . . .	14

## Background


Today we will begin our exploration of some important machine learning methods, namely **clustering** and **dimensionality reduction**.

Let's make up some input data for clustering where we know what the natural "clusters" are.

The function `rnorm()` can be useful here.


```
hist( rnorm(5000, mean=10) )
```

### Histogram of rnorm(5000, mean = 10)



Q. Generate 30 random numbers centered at +3 and another 30 centered at -3

```
tmp <- c(rnorm(30, mean=3),  
         rnorm(30, mean=-3))  
x <- cbind(x=tmp, y=rev(tmp))  
plot(x)
```



## K-means clustering

The main function in “base R” for K-means clustering is called `kmeans()`:

```
km <- kmeans(x, centers=2)  
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.665887	2.690928
2	2.690928	-2.665887

Clustering vector:

Within cluster sum of squares by cluster:

```
[1] 60.8784 60.8784
```

(between SS / total SS = 87.6 %)

## Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

Q. What component of the results object details the cluster size?

km\$size

[1] 30 30

Q. What component of the results object details the cluster centers?

km\$centers


	x	y
1	-2.665887	2.690928
2	2.690928	-2.665887

Q. What component of the results object details the cluster membership vector (i.e. our main result of which points lie in which cluster)?

km\$cluster


Q. Plot our clustering results with points colored by cluster and also add the cluster centers as new points colored blue?

```
plot(x, col= c(1,2))  
points(km$centers, col="blue", pch=15)
```



Q. Run `kmeans()` again and this time produce 4 clusters (and call your result object `k4`) and make a results figure like above?

```
k4 <- kmeans(x, centers=4)
plot(x, col=k4$cluster)
points(k4$center, col="blue", pch=15)
```



The metric

```
km$tot.withinss
```

```
[1] 121.7568
```

```
k4$tot.withinss
```


```
[1] 77.62604
```

Q. Let's try different number of k (centers) from 1 to 30 and see what the best results is?

```
i <- 1
ans <- NULL
for (i in 1:30) {
  ans <- c(ans, kmeans(x, centers=i)$tot.withinss)
}
ans
```

```
[1] 982.620801 121.756810 95.637217 85.734450 57.788111 46.058599
[7] 38.812114 45.366701 50.073939 26.390376 22.255789 19.154714
[13] 15.985639 19.842222 15.568101 13.792783 14.061086 8.101669
[19] 9.740967 10.826640 13.631115 5.962910 4.558157 4.943509
[25] 4.647703 5.048145 4.011763 3.880538 5.069867 4.945483
```

```
plot(ans, typ="o")
```




**Key-Point:** K-means will impose a clustering structure on your data even if it is not there - it will always give you the answer you asked for even if that answer is silly!

## Hierarchical Clustering

The main function for Hierarchical clustering is called `hclust()`. Unlike `kmeans()` (which does all the work for you) you can't just pass `hclust()` our raw input data. It needs a "distance matrix" like the one returned from the `dist()` function.

```
d <- dist(x)
hc <- hclust(d)
plot(hc)
```


## Cluster Dendrogram



```
d  
hclust (*, "complete")
```

To extract our cluster membership vector from a `hclust()` result object we have to “cut” our tree at a given height to yield separate “groups”/“branches”.

```
plot(hc)  
abline(h=8, col="red", lty=2)
```



To do this we used the `cutree()` function on our `hclust()` object:

```
grps <- cutree(hc, h=8)  
grps
```

```
table(grps, km$cluster)
```

grps	1	2
1	0	30
2	30	0

## PCA of UK food data

Import the dataset of food consumption in the UK:

```
url <- "https://tinyurl.com/UK-foods"  
x <-read.csv(url)  
x
```

		X	England	Wales	Scotland	N.Ireland
1	Cheese		105	103	103	66
2	Carcass_meat		245	227	242	267
3	Other_meat		685	803	750	586
4	Fish		147	160	122	93
5	Fats_and_oils		193	235	184	209
6	Sugars		156	175	147	139
7	Fresh_potatoes		720	874	566	1033
8	Fresh_Veg		253	265	171	143
9	Other_Veg		488	570	418	355
10	Processed_potatoes		198	203	220	187
11	Processed_Veg		360	365	337	334
12	Fresh_fruit		1102	1137	957	674
13	Cereals		1472	1582	1462	1494
14	Beverages		57	73	53	47
15	Soft_drinks		1374	1256	1572	1506
16	Alcoholic_drinks		375	475	458	135
17	Confectionery		54	64	62	41

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17 5
```

One solution to set the row names is to do it by hand...

```
rownames(x) <- x[,1]
```

To remove the first column I can use the minus index trick

```
x <- x[,-1]
```

A better way to do this is to set the row names to the first column with `read.csv()`

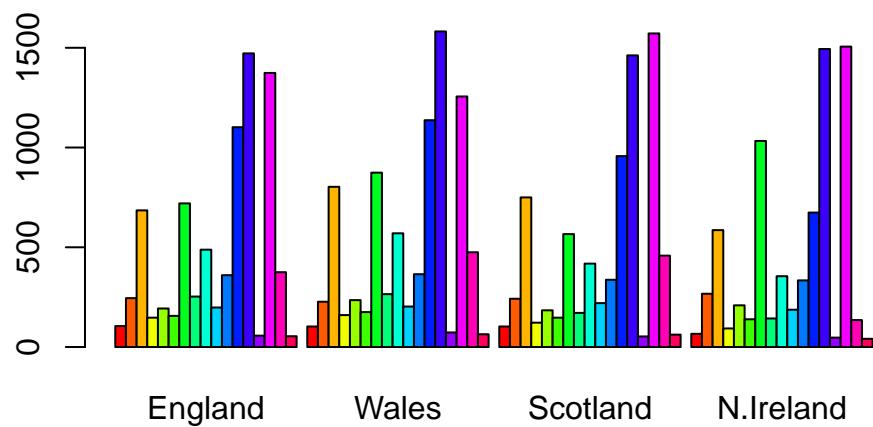
```
x <- read.csv(url, row.names=1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

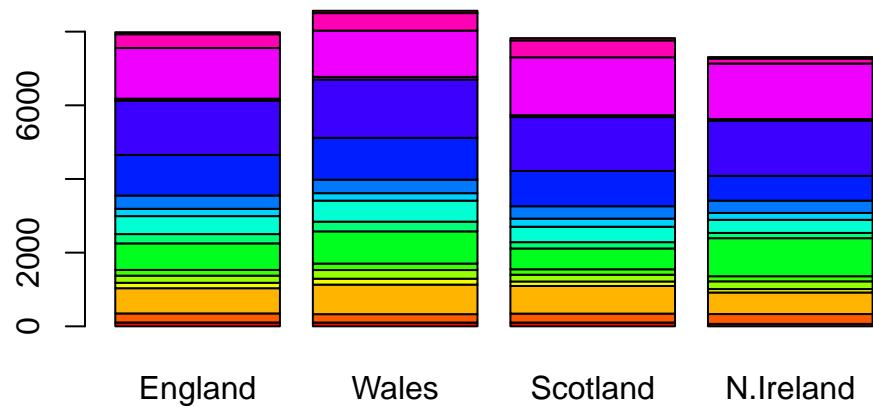
### Spotting major differences and trends

It is difficult even in this wee 17D dataset...

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```

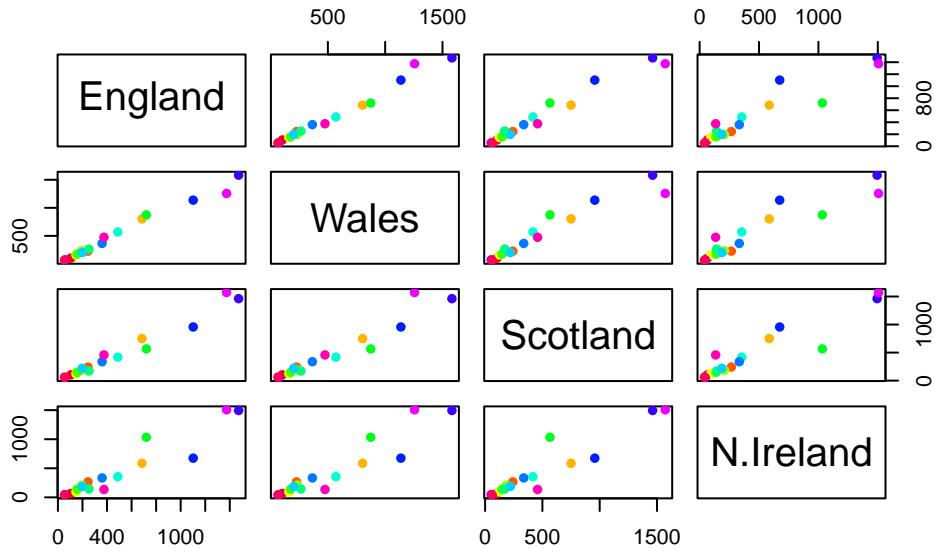


```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```

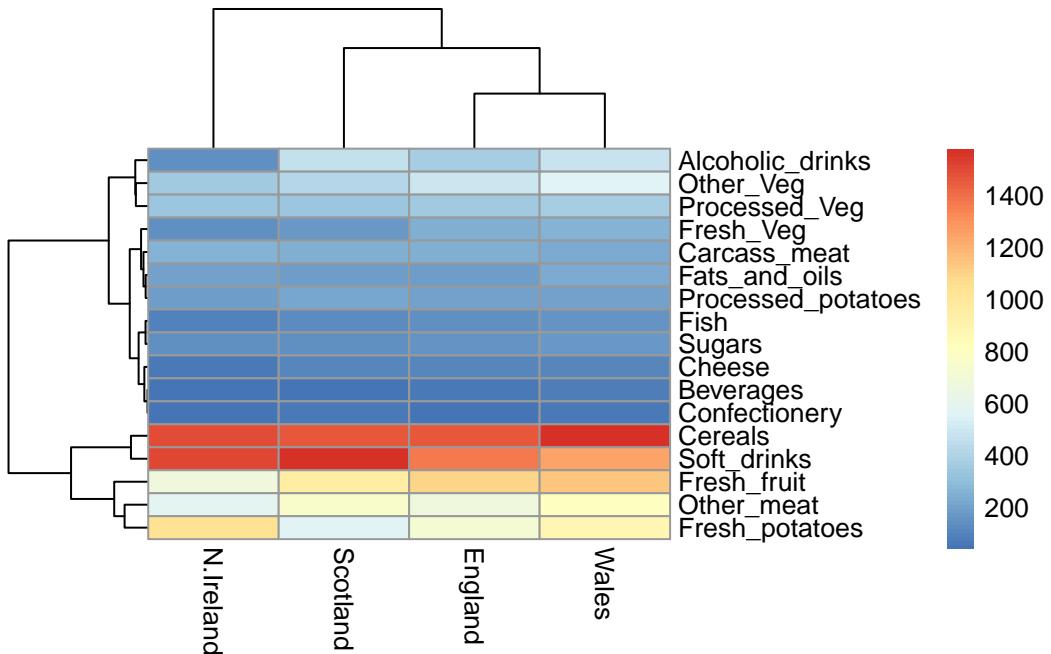


## Paris Plots and heat maps

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



```
library(pheatmap)
pheatmap( as.matrix(x) )
```



## PCA to the rescue

The main PCA function in “base R” is called `prcomp()`. This function wants the transpose of our food data as input (i.e. the foods as columns and the countries as rows).

```
pca <- prcomp( t(x) )
```

```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation"   "center"     "scale"      "x"


$class
[1] "prcomp"
```

To make one of main PCA result figures we turn to `pca$x` the scores along our new PCs. This is called “PC plot” or “score plot” or “oriendation plot”...

```
my_cols<- c("orange", "red", "blue", "darkgreen")
```

```
library(ggplot2)

ggplot(pca$x)+
  aes(PC1, PC2)+
  geom_point(col= my_cols)
```



The second major result figure is called a “loadings plot” of “variable contributions plot” or “weight plot”

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.694538519
Carcass_meat	0.047927628	0.013915823	0.06367111	0.489884628
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.279023718
Fish	-0.084414983	-0.050754947	0.03906481	-0.008483145
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.076097502

Sugars	-0.037620983	-0.043021699	-0.03605745	0.034101334
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	-0.090972715
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	-0.039901917
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.016719075
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	0.030125166
Processed_Veg	-0.036488269	-0.045451802	0.05289191	-0.013969507
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.184072217
Cereals	-0.047702858	-0.212599678	-0.35884921	0.191926714
Beverages	-0.026187756	-0.030560542	-0.04135860	0.004831876
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.103508492
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.316290619
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001847469

```
ggplot(pca$rotation)+  
  aes(PC1, rownames(pca$rotation))+  
  geom_col()
```

