

Class 08: Breast Cancer Mini Project

Melanie Alonzo (PID: A17375327)

Table of contents

Background	1
Data Import	1
Principal Component Analysis (PCA)	5
Communicating PCA results	10
4. Hierarchical Clustering	10
Combining methods	11
Prediction	13

Background

In today's class we will be employing all the R techniques for data analysis that we have learned thus far - including the machine learning methods of clustering and PCA - to analyze real breast cancer biopsy data.

Data Import

The data is in CSV format:

```
fna.data <- "WisconsinCancer.csv"  
wisc.df <- read.csv(fna.data, row.names=1)
```

We peek at the data

```
head(wisc.df, 3)
```

```

diagnosis radius_mean texture_mean perimeter_mean area_mean
842302      M      17.99      10.38      122.8      1001
842517      M      20.57      17.77      132.9      1326
84300903      M      19.69      21.25      130.0      1203
smoothness_mean compactness_mean concavity_mean concave.points_mean
842302      0.11840      0.27760      0.3001      0.14710
842517      0.08474      0.07864      0.0869      0.07017
84300903      0.10960      0.15990      0.1974      0.12790
symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
842302      0.2419      0.07871      1.0950      0.9053      8.589
842517      0.1812      0.05667      0.5435      0.7339      3.398
84300903      0.2069      0.05999      0.7456      0.7869      4.585
area_se smoothness_se compactness_se concavity_se concave.points_se
842302      153.40      0.006399      0.04904      0.05373      0.01587
842517      74.08      0.005225      0.01308      0.01860      0.01340
84300903      94.03      0.006150      0.04006      0.03832      0.02058
symmetry_se fractal_dimension_se radius_worst texture_worst
842302      0.03003      0.006193      25.38      17.33
842517      0.01389      0.003532      24.99      23.41
84300903      0.02250      0.004571      23.57      25.53
perimeter_worst area_worst smoothness_worst compactness_worst
842302      184.6      2019      0.1622      0.6656
842517      158.8      1956      0.1238      0.1866
84300903      152.5      1709      0.1444      0.4245
concavity_worst concave.points_worst symmetry_worst
842302      0.7119      0.2654      0.4601
842517      0.2416      0.1860      0.2750
84300903      0.4504      0.2430      0.3613
fractal_dimension_worst
842302      0.11890
842517      0.08902
84300903      0.08758

```

Q1. How many observations are in this dataset?

```
nrow(wisc.df)
```

```
[1] 569
```

Q2. How many of the observations have a malignant diagnosis?

```
sum (wisc.df$diagnosis == "M" )
```

```
[1] 212
```

```
table (wisc.df$diagnosis)
```

```
B      M  
357 212
```

Q3. How many variables/features in the data are suffixed with _mean?

```
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
84358402	0.10030	0.13280	0.1980	0.10430	
843786	0.12780	0.17000	0.1578	0.08089	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
84300903	0.2069		0.05999	0.7456	0.7869
84348301	0.2597		0.09744	0.4956	1.1560
84358402	0.1809		0.05883	0.7572	0.7813
843786	0.2087		0.07613	0.3345	0.8902
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
84348301	27.23	0.009110	0.07458	0.05661	0.01867

84358402	94.44	0.011490	0.02461	0.05688	0.01885
843786	27.19	0.007510	0.03345	0.03672	0.01137
		symmetry_se	fractal_dimension_se	radius_worst	texture_worst
842302	0.03003		0.006193	25.38	17.33
842517	0.01389		0.003532	24.99	23.41
84300903	0.02250		0.004571	23.57	25.53
84348301	0.05963		0.009208	14.91	26.50
84358402	0.01756		0.005115	22.54	16.67
843786	0.02165		0.005082	15.47	23.75
		perimeter_worst	area_worst	smoothness_worst	compactness_worst
842302	184.60	2019.0		0.1622	0.6656
842517	158.80	1956.0		0.1238	0.1866
84300903	152.50	1709.0		0.1444	0.4245
84348301	98.87	567.7		0.2098	0.8663
84358402	152.20	1575.0		0.1374	0.2050
843786	103.40	741.6		0.1791	0.5249
		concavity_worst	concave.points_worst	symmetry_worst	
842302	0.7119		0.2654	0.4601	
842517	0.2416		0.1860	0.2750	
84300903	0.4504		0.2430	0.3613	
84348301	0.6869		0.2575	0.6638	
84358402	0.4000		0.1625	0.2364	
843786	0.5355		0.1741	0.3985	
		fractal_dimension_worst			
842302		0.11890			
842517		0.08902			
84300903		0.08758			
84348301		0.17300			
84358402		0.07678			
843786		0.12440			

```
length(grep("_mean", colnames(wisc.df)))
```

```
[1] 10
```

```
colnames(wisc.df)
```

```
[1] "diagnosis"                  "radius_mean"
[3] "texture_mean"               "perimeter_mean"
[5] "area_mean"                  "smoothness_mean"
[7] "compactness_mean"            "concavity_mean"
```

```
[9] "concave.points_mean"      "symmetry_mean"
[11] "fractal_dimension_mean"   "radius_se"
[13] "texture_se"              "perimeter_se"
[15] "area_se"                 "smoothness_se"
[17] "compactness_se"          "concavity_se"
[19] "concave.points_se"       "symmetry_se"
[21] "fractal_dimension_se"    "radius_worst"
[23] "texture_worst"           "perimeter_worst"
[25] "area_worst"              "smoothness_worst"
[27] "compactness_worst"       "concavity_worst"
[29] "concave.points_worst"    "symmetry_worst"
[31] "fractal_dimension_worst"
```

We need to remove the `diagnosis` column before we do any further analysis of this dataset - we dont want to pass this to PCA etc. We will save it as a separate wee vector that we can use later to compare our findings to those of experts.

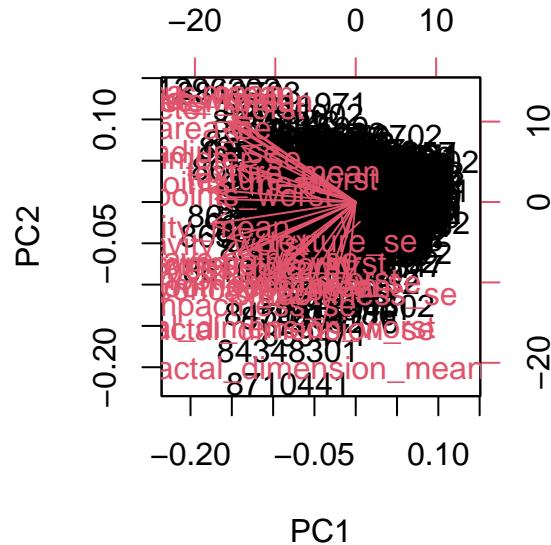
```
wisc.data <- wisc.df[,-1]
diagnosis <- wisc.df$diagnosis
```

Principal Component Analysis (PCA)

The main function in base R is called `prcomp()` we will use the optional argument `scale=TRUE` here as the data columns/ features/ dimensions are on very different scales in the original data set.

```
wisc.pr <- prcomp(wisc.data, scale=T)
```

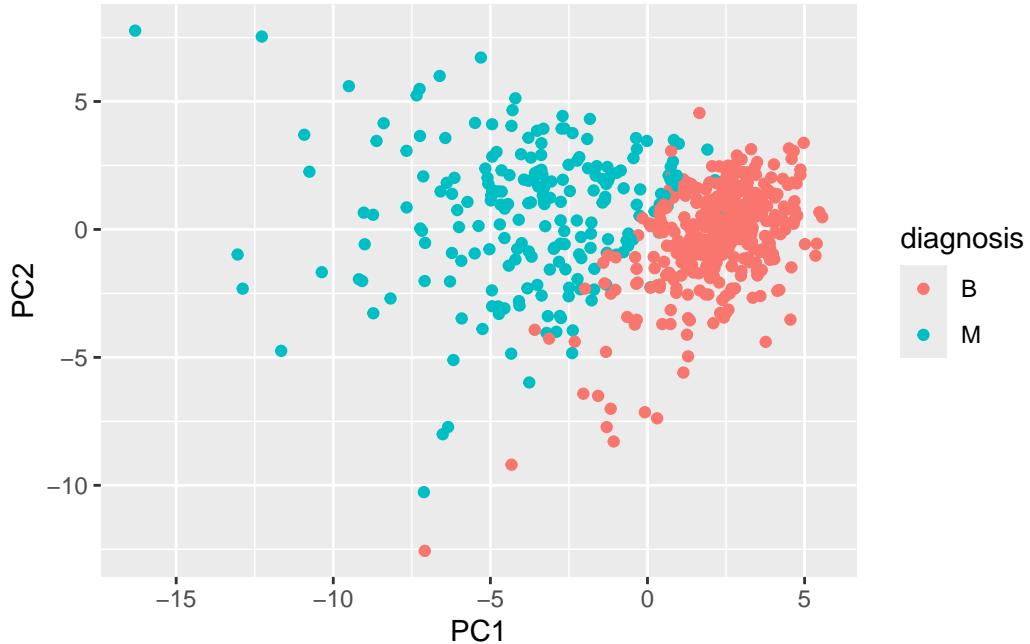
```
biplot(wisc.pr)
```



```
attributes(wisc.pr)
```

```
$names
[1] "sdev"      "rotation"   "center"    "scale"     "x"
$class
[1] "prcomp"
```

```
library(ggplot2)
ggplot(wisc.pr$x) +
  aes(PC1, PC2, col= diagnosis) +
  geom_point()
```

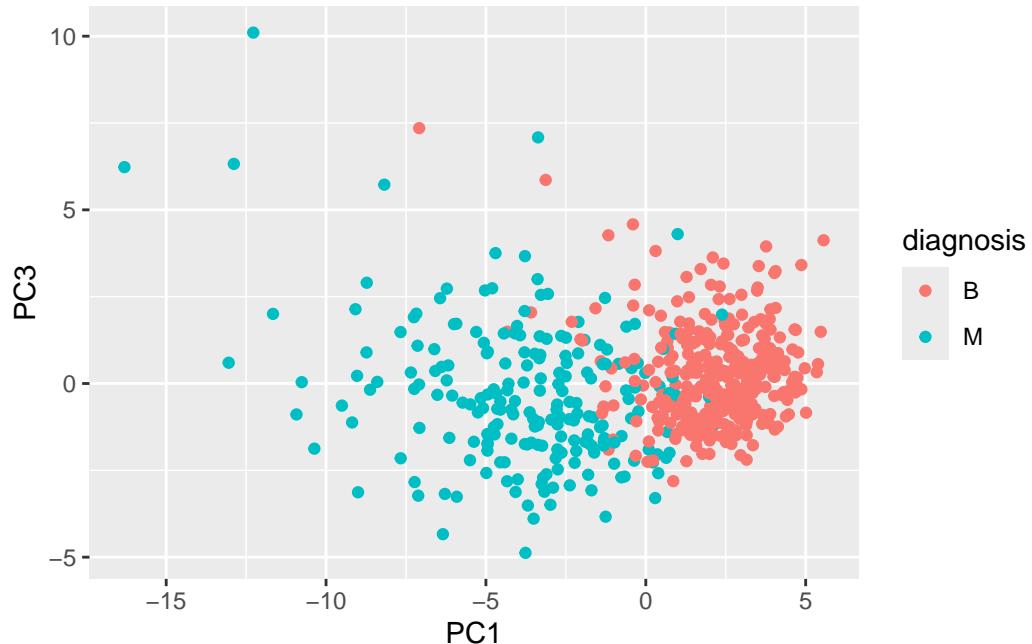


```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

```
ggplot(wisc.pr$x) +  
  aes(PC1, PC3, col=diagnosis) +  
  geom_point()
```

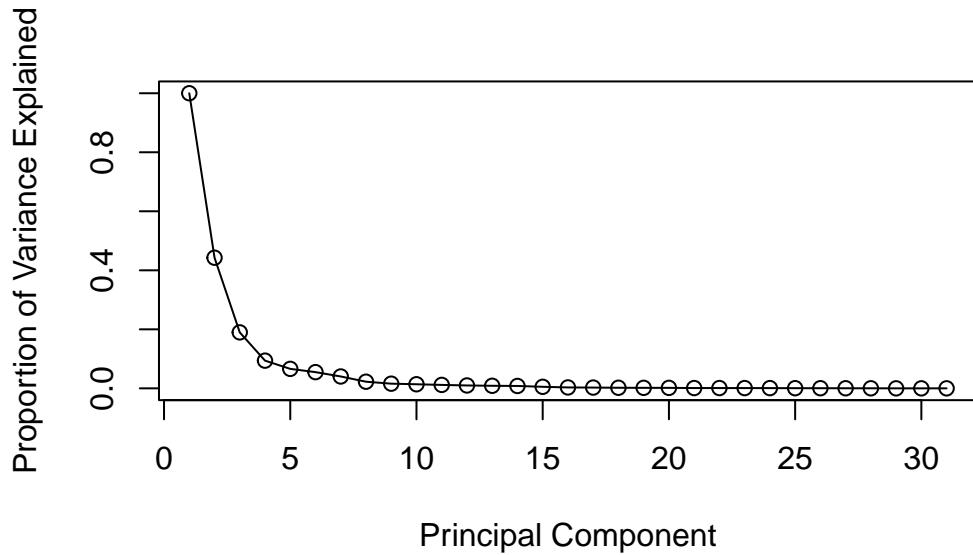


```
pr.var <- wisc.pr$sdev^2  
head(pr.var)
```

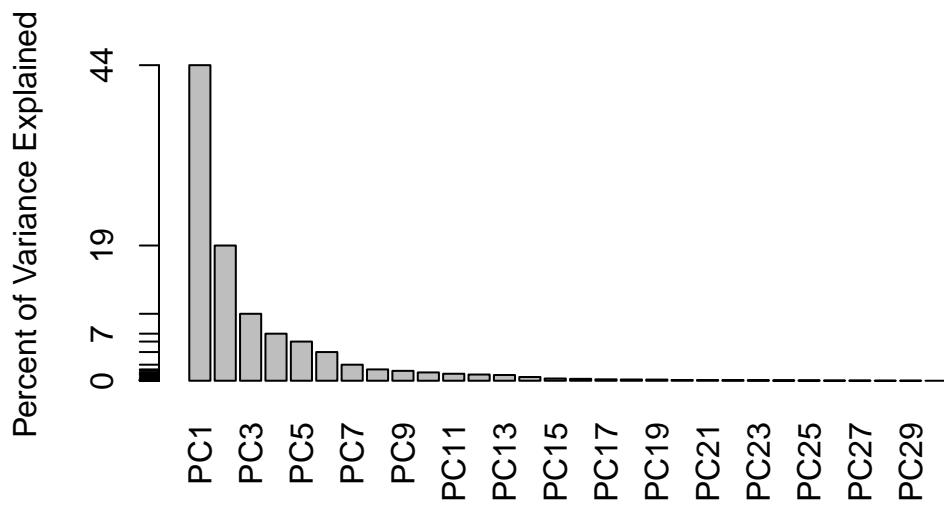
```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

```
pve <- pr.var / sum(pr.var)
```

```
plot(c(1,pve), xlab = "Principal Component",  
     ylab = "Proportion of Variance Explained",  
     ylim = c(0, 1), type = "o")
```



```
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

```
wisc.pr$rotation ["concave.points_mean", "PC1"]
```

```
[1] -0.2608538
```

4. Hierarchical Clustering

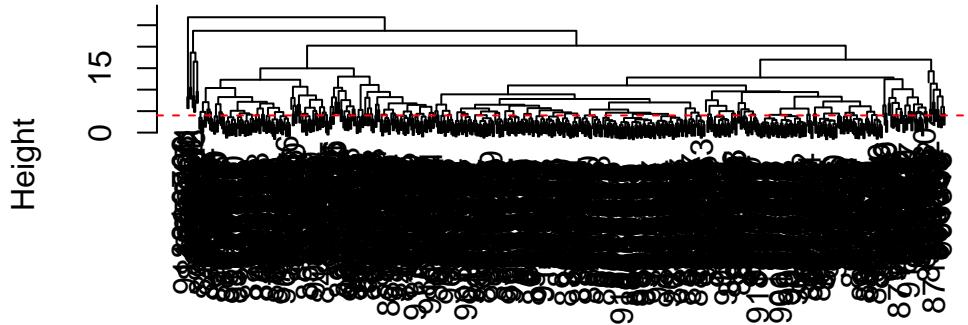
The goal of this section is to do hierarchical clustering of the **original data** to see if there is any obvious grouping into malignant and benign clusters.

In short, these results are not good!

First we will scale our `wisc.data` then calculate a distance matrix, then pass to `hclust()`:

```
wisc.dist <- dist( scale(wisc.data))
wisc.hclust<- hclust( wisc.dist)
plot(wisc.hclust)
abline(h=4, col="red", lty=2)
```

Cluster Dendrogram



```
wisc.dist  
hclust (*, "complete")
```

```
wisc.hclust.clusters <- cutree(wisc.hclust,k=2)  
table(wisc.hclust.clusters)
```

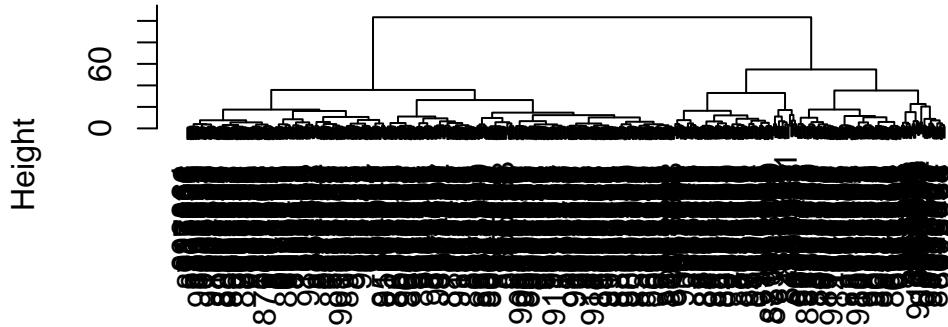
```
wisc.hclust.clusters  
1 2  
567 2
```

Combining methods

The idea here is that I can take my new variables (i.e. the score on the PCs `wisc.pr$x`) that are better descriptors of the data-set than the original features (i.e. the 30 columns in `wisc.data`) and use these as a basis for clustering.

```
pc.dist <- dist(wisc.pr$x[,1:3])  
wisc.pr.hclust <- hclust(pc.dist, method= "ward.D2")  
plot(wisc.pr.hclust)
```

Cluster Dendrogram



```
pc.dist  
hclust (*, "ward.D2")
```

```
grps <- cutree(wisc.pr.hclust, k=2)  
table(grps)
```

```
grps  
1 2  
203 366
```

```
table(diagnosis)
```

```
diagnosis  
B M  
357 212
```

I can now run `table()` with both my clustering `grps` and the expert diagnosis

```
table(grps, diagnosis)
```

```
diagnosis  
grps B M  
1 24 179  
2 333 33
```

Our cluster “1” has 179 “M” diagnosis or cluster “2” has 333 “B” diagnosis

- 179 TP
- 24 FP
- 333 TN
- 33 FN

Sensitivity: $TP/(TP+FN)$

```
179/(179+33)
```

```
[1] 0.8443396
```

Specificity: $TN/(TN+FP)$

```
333/(333+24)
```

```
[1] 0.9327731
```

Prediction

We can use our PCA model for prediction of new un-seen cases:

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	

```
[2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
      PC27        PC28        PC29        PC30
[1,]  0.220199544 -0.02946023 -0.015620933  0.005269029
[2,] -0.001134152  0.09638361  0.002795349 -0.019015820
```

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```

